

Radar Search Pipeline TO BE 0.7.6.4.18.1.1

Review design for a planned Radar search pipeline change

Status: Implemented design input

Product area: Radar candidate discovery pipeline

Slice: 0.7.6.4.18.1.1 Candidate discovery recall-first upstream semantics and benchmark target protection

Last updated: 2026-07-06

Generated PDF: docs/radar/to-be/RADAR_SEARCH_PIPELINE_TO_BE_0.7.6.4.18.1.1.pdf

Goal

Candidate discovery must behave as a recall-first upstream finder after the signal-monitoring split. Source-backed upstream leads must not collapse into a flat Monitor result only because signal monitoring has not run. Downstream product acceptance stays strict and separate.

Intended Behavior

1. Candidate discovery keeps source-backed upstream leads when there is valid retrieved-source, official-domain, or registry identity evidence.
2. upstream_discovery_outcome explains what discovery found: confirmed_upstream_lead, review_needed_upstream_lead, retained_upstream_lead, or rejected_noise.
3. product_acceptance_status explains whether the entity is a strict product candidate, requires review, or is not product accepted.
4. In normal signal_execution_mode="handoff", missing signal observations do not create Monitor, signal_requires_human_review, or not_observed.
5. Official high-trust domains can promote source-backed relation and industrial coverage evidence when the candidate name or alias appears in product-safe source diagnostics.
6. Concrete registry identity evidence, such as INN, OGRN, or legal name, is retained as upstream evidence and is not weak by construction.
7. benchmark_live protects explicit baseline targets before optional exploration by carrying lane minimums, completion slots, reserve budgets, and benchmark metadata through dedupe, selection, scheduling, and reporting.

Reporting Contract

Evaluation reports separate upstream recall from strict product precision:

- product_candidate_count counts only strict product-accepted candidates;
- retained_upstream_lead_count counts retained upstream material;
- confirmed_upstream_lead_count and review_needed_upstream_lead_count explain the quality of retained upstream leads;

- `benchmark_target_funnel` gives every baseline target a path state such as generated, selected, admitted, executed, source found, projected, or a path-level miss reason.

Guardrails

- No signal-monitoring live runtime is added in this slice.
- No SIBUR production hardcode is allowed outside benchmark fixtures/hints.
- Upstream false positives are acceptable as retained/review-needed leads, but must not become accepted product accounts without strict evidence.
- Every remaining false negative must have a path-level reason instead of a generic budget symptom.

Validation Plan

- Unit tests cover official-domain admission, open-web retention, registry identity retention, handoff normalization, benchmark target metadata dedupe, and benchmark-present source projection.
- Regression gates cover search expansion, benchmark profile generation, evaluation reports, live candidate discovery, adaptive execution, API/jobs, signal-monitoring isolation, documentation, roadmap, and static diff checks.