

Radar Candidate Discovery TO BE 0.7.6.4.18.1.4

Review design for a planned Radar search pipeline change

1. Slice And Decision Context

Slice 0.7.6.4.18.1.4 corrects the candidate-discovery result surface after the recovery and reconciliation repairs. The current pipeline can retain broad upstream legal targets, but the user-facing candidate list is still too narrow: the benchmark smoke can match eight legal baseline targets while only three rows are visible as public/product candidates.

The deliberate decision is to split the visible candidate surface from strict product acceptance:

- accepted product candidates are strict account candidates;
- review-needed candidates are user-visible legal leads that need human or downstream acceptance work;
- universe-only diagnostics remain available for non-account entities, noisy rows, gaps, and traceable explanations.

Signal monitoring remains out of scope. Candidate discovery still ends with a handoff snapshot and must not create searched-negative signal observations.

2. AS IS Problem Statement

Implemented behavior after 0.7.6.4.18.1.3:

- upstream leads no longer disappear without a reconciliation reason;
- dossier and evaluation expose candidate_discovery_reconciliation and product_acceptance_ledger;
- the SIBUR benchmark smoke can report high strict recall over candidate_universe;
- public candidate rows still represent only the strict accepted subset.

This leaves a product gap: users see three public/product candidates, while five additional source-backed legal baseline targets are buried in diagnostic surfaces. The result is technically explainable but still product-poor.

3. Intended Pipeline Behavior

Candidate discovery now produces three explicit surfaces:

1. candidate_universe: broad upstream truth, including legal entities, sites, branches, ambiguous rows, gaps, and diagnostics.
2. user_visible_candidates: the user-facing candidate surface. It contains: accepted product candidates and review-needed legal candidates.
3. strict product candidates: the accepted subset used for product precision.

Public candidate rows must carry a surface status:

- accepted_product_candidate;
- review_needed_candidate;
- universe_only_diagnostic;
- not_promoted.

Review-needed legal candidates are visible, but they do not inflate precision and do not become accepted product candidates.

4. Roles Changed

Candidate projection/finalization

New responsibility	Build user-visible accepted and review-needed candidate rows from source-backed legal universe entries.
Still does not own	Provider calls, signal monitoring, or product precision rules outside candidate discovery.

Outcome reconciliation

New responsibility	Count accepted, review-needed, visible, universe-only, and not-promoted rows consistently.
Still does not own	Retrieval, extraction, or benchmark scoring.

API/dossier mappers

New responsibility	Expose the same user-visible candidate rows and candidate universe with surface status.
Still does not own	Mutating run behavior.

Evaluation

New responsibility	Report visible legal baseline coverage separately from strict product precision.
Still does not own	Provider calls or hiding quality failures.

Search expansion/target funnel

New responsibility	Give protected targets such as Poliom a specific bounded reason when they are generated but not selected.
Still does not own	Unbounded retry or hardcoded production selection.

5. Context Passed Between Roles

Finalization passes product-safe fields only:

- candidate name, entity type, source refs, benchmark id, aliases;
- upstream discovery outcome and confidence;
- product acceptance status and reason;
- public projection status and reason;
- candidate surface status and reason;
- target-funnel path reason for protected benchmark targets.

Raw prompts, hidden reasoning, credentials, headers, and raw provider dumps are not allowed in the visible surface or evaluation report.

6. Source, Budget, And Checkpoint Semantics

The slice does not add broad fallback or unbounded budget.

Source-backed legal entities may become review-needed visible candidates when:

- they are legal entities;
- they have source refs or structured registry identity;

- they match protected benchmark legal targets or satisfy deterministic

source-backed admission rules.

Budget and checkpoint semantics remain bounded. A generated-but-not-selected protected target must receive a specific bounded reason such as selection cap, reserve cap, admission cap, or execution cap. Generic disappearance is not allowed.

7. Dossier, Trace, And Evaluation Visibility

The dossier and candidates endpoint expose:

- candidate_surface_status;
- accepted vs review-needed candidate counts;
- visible legal baseline count where evaluation context exists;
- strict product candidate count;
- product acceptance reason;
- surface promotion reason.

The evaluation report exposes:

- legal_baseline_visible_count;
- accepted_product_candidate_count;
- review_needed_candidate_count;
- visible_candidate_count;
- strict_product_candidate_count;
- unchanged strict_recall, review_recall, precision;
- target-funnel reasons for all baseline targets.

8. Changed Flow

9. Test Plan

Fast tests must prove:

- source-backed legal universe entries can become review-needed visible rows;
- accepted and review-needed candidates are counted separately;
- review-needed visible candidates do not count as accepted product candidates;
- the SIBUR-style fixture reaches at least eight visible legal baseline targets

and at least three accepted product candidates;

- protected targets cannot regress to present_not_projected;
- not_observed is not created in handoff mode.

Final gate:

1. rebuild Docker with docker compose up -d --build;
2. run SIBUR benchmark_smoke;
3. evaluate the run through API;
4. close the slice only when the DoD metrics pass.

10. Acceptance Criteria

- legal_baseline_visible_count >= 8;
- accepted_product_candidate_count >= 3;
- review_needed_candidate_count >= 5, or every missing legal target has a specific target-funnel and ledger reason;
- unexplained_drop_count == 0;
- present_not_projected_count == 0;
- Poliom is visible/projected or has a specific bounded selection/cap reason;
- handoff signal rows remain not-searched/pending/limited.

11. Explicit Out Of Scope

- no signal-monitoring live runtime;
- no quality claim from one live run;
- no SIBUR-specific production hardcode outside benchmark fixtures/hints;
- no unbounded budget increase;
- no lowering strict product precision.

12. Open Questions

- Whether review-needed legal candidates should appear in the same UI table as accepted candidates or in a separate grouped section is a frontend question for a later UI slice. The backend contract must expose both either way.